

# INT104 ARTIFICIAL INTELLIGENCE

## L9- Unsupervised Learning I Clustering

Fang Kang

[Fang.kang@xjtlu.edu.cn](mailto:Fang.kang@xjtlu.edu.cn)



Xi'an Jiaotong-Liverpool University

西交利物浦大學




# CONTENT


- K-Means
  - ◆ K-means clustering
  - ◆ Centroid initialization methods
  - ◆ Parameters and Evaluation
- Hierarchical Clustering
- DBSCAN




# Supervised vs. unsupervised



Supervised



Unsupervised





Semi-supervised




# Supervised learning

- The correct labels for each training example are known



# Unsupervised learning

Labels are unknown → Need to automatically discover the ***clustering*** pattern and structure in data



# Real world clustering example

- Clustering on text documents

Google News

Top stories  
For you  
Favorites  
Saved searches

U.S.

World

Local

Business

Technology

Entertainment

Sports

Science

Health


- Image segmentation



- Streaming services often use clustering analysis to identify viewers who have similar behavior.



# K-means clustering algorithm




Term first used by James MacQueen,  
1967

Algorithm proposed by Stuart Lloyd,  
1957, published in 1982.



# K-means clustering algorithm



Goal: Assign all data points to 2 clusters


**Step 1:** Pick 2 *random* initial cluster centroids

**Step 2:** Paint the data points that are closer to red centroid **red**, and those closer to blue centroid **blue**





# K-means clustering algorithm



Goal: Assign all data points to 2 clusters

**Step 1:** Pick 2 *random* initial cluster centroids

**Step 2:** Paint the data points that are closer to red centroid **red**, and those closer to blue centroid **blue**


**Step 3:** Update the positions of centroids

**Red** centroid := average of current red points

**Blue** centroid := average of current blue points



# K-means clustering algorithm



Goal: Assign all data points to 2 clusters

**Step 1:** Pick 2 *random* initial cluster centroids

**Step 2:** Paint the data points that are closer to red centroid **red**, and those closer to blue centroid **blue**

**Step 3:** Update the positions of centroids


**Red** centroid := average of current red points

**Blue** centroid := average of current blue points

Repeat



# K-means clustering algorithm



Goal: Assign all data points to 2 clusters

**Step 1:** Pick 2 *random* initial cluster centroids

**Step 2:** Paint the data points that are closer to red centroid **red**, and those closer to blue centroid **blue**

**Step 3:** Update the positions of centroids


**Red** centroid := average of current red points

**Blue** centroid := average of current blue points

Repeat



# K-means clustering algorithm



Goal: Assign all data points to 2 clusters

**Step 1:** Pick 2 *random* initial cluster centroids

**Step 2:** Paint the data points that are closer to red centroid **red**, and those closer to blue centroid **blue**

**Step 3:** Update the positions of centroids

**Red** centroid := average of current red points

**Blue** centroid := average of current blue points

Repeat

Until no more points need to be repainted, i.e., the centroids no longer change

Clustering is done



# K-means formal definition

Given a dataset  $\{x^{(1)}, \dots, x^{(m)}\}$ ,  $x^{(i)} \in \mathbb{R}^n$ , and want to group the data into  $k$  clusters

1. Initialize  $k$  **cluster centroids**  $\mu_1, \mu_2, \dots, \mu_k \in \mathbb{R}^n$  randomly

2. Repeat until convergence: {

For  $i = 1, \dots, m$ :

$c^{(i)} := \arg \min_j \|x^{(i)} - \mu_j\|^2 \longrightarrow$  Assign  $x^{(i)}$  to the closest cluster  $j$

Usually, Euclidean distance is the best option

For  $j = 1, \dots, k$ :

$\mu_j := \frac{\sum_{i=1}^m 1\{c^{(i)}=j\}x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)}=j\}} \longrightarrow$  Update centroid  $\mu_j$  with mean of all within-cluster data points

}

From a machine learning perspective, K-means minimize the cost function:

$$J(c, \mu) = \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2 \quad \text{guaranteed to converge}$$

The **City block distance** between two points,  $a$  and  $b$ , with  $k$  dimensions is calculated as:




$$\sum_{j=1}^k |a_j - b_j|$$



# How to initialize centroids?

Randomly pick  $k$  data points as the initial centroids

Sometimes it leads to different clustering results



Solutions:



- Run multiple times with different initializations and evaluate
- K-means++ (Arthur & Vassilvitskii, 2007)



# Better initialization with $K$ -means++


Arthur & Vassilvitskii, 2007

1. Choose one centroid uniformly at random from data points.
2. For each  $x^{(i)}$ , compute  $D(x^{(i)})$ , the distance between  $x^{(i)}$  and the nearest centroid that has already been chosen.
3. Choose one new data point at random as a new centroid, where the probability of choosing point  $x^{(i)}$  is **proportional** to  $D(x^{(i)})$ .
4. Repeat until  $k$  centroids have been chosen.




# How to choose $K$ ?

By intuitive observation



Better use one cluster



Better use two clusters

Is there a systematic evaluation?





# Parameters and Evaluation

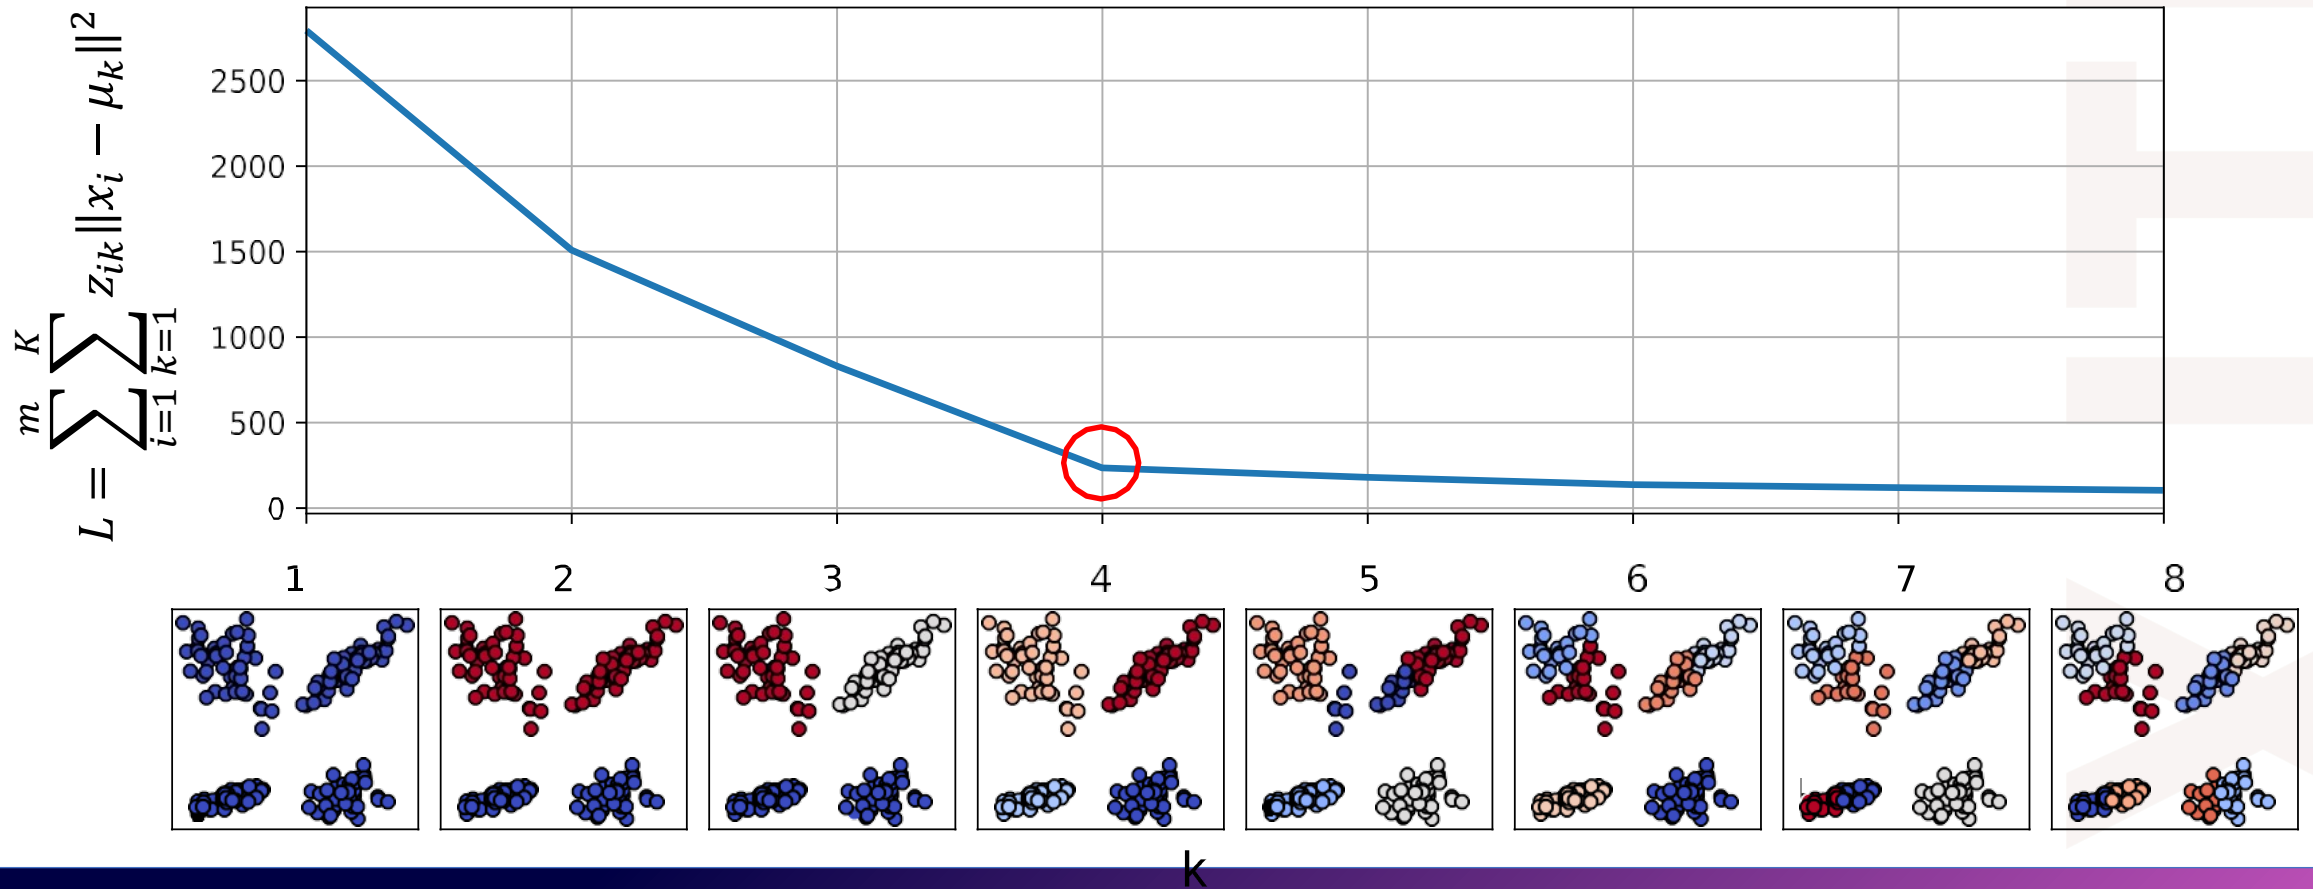
- The number of clusters  $k$  is a hyperparameter. How do we find a good  $k$ ?
  1. Elbow method:
    - Start with a small  $k$  value and increase it until adding another cluster does not result in a much lower distortion value
    - In other words, the new cluster does not explain so much the variance in data
  2. Silhouette Coefficient:
    - A measure of how tight each cluster is and how far apart clusters are from each other
    - Choose a value of  $k$  that results in clustering with a large silhouette coefficient

Silhouette /ˌsɪlʊˈet/: contour, outline



# The Elbow Method

- Choose  $k$  such that adding another cluster will not explain the variance in data by much (i.e. does not give a much lower distortion value)



# The Silhouette Coefficient

A good clustering algorithm:  $\left\{ \begin{array}{l} \text{high similarity **within** cluster} \\ \text{low similarity **between** cluster} \end{array} \right.$

- **silhouette coefficient** Measures the tightness of clusters and separation between clusters:

$$S = \frac{1}{m} \sum_{i=1}^m s(x_i)$$
$$s(x_i) = \frac{b(x_i) - a(x_i)}{\max\{a(x_i), b(x_i)\}}$$

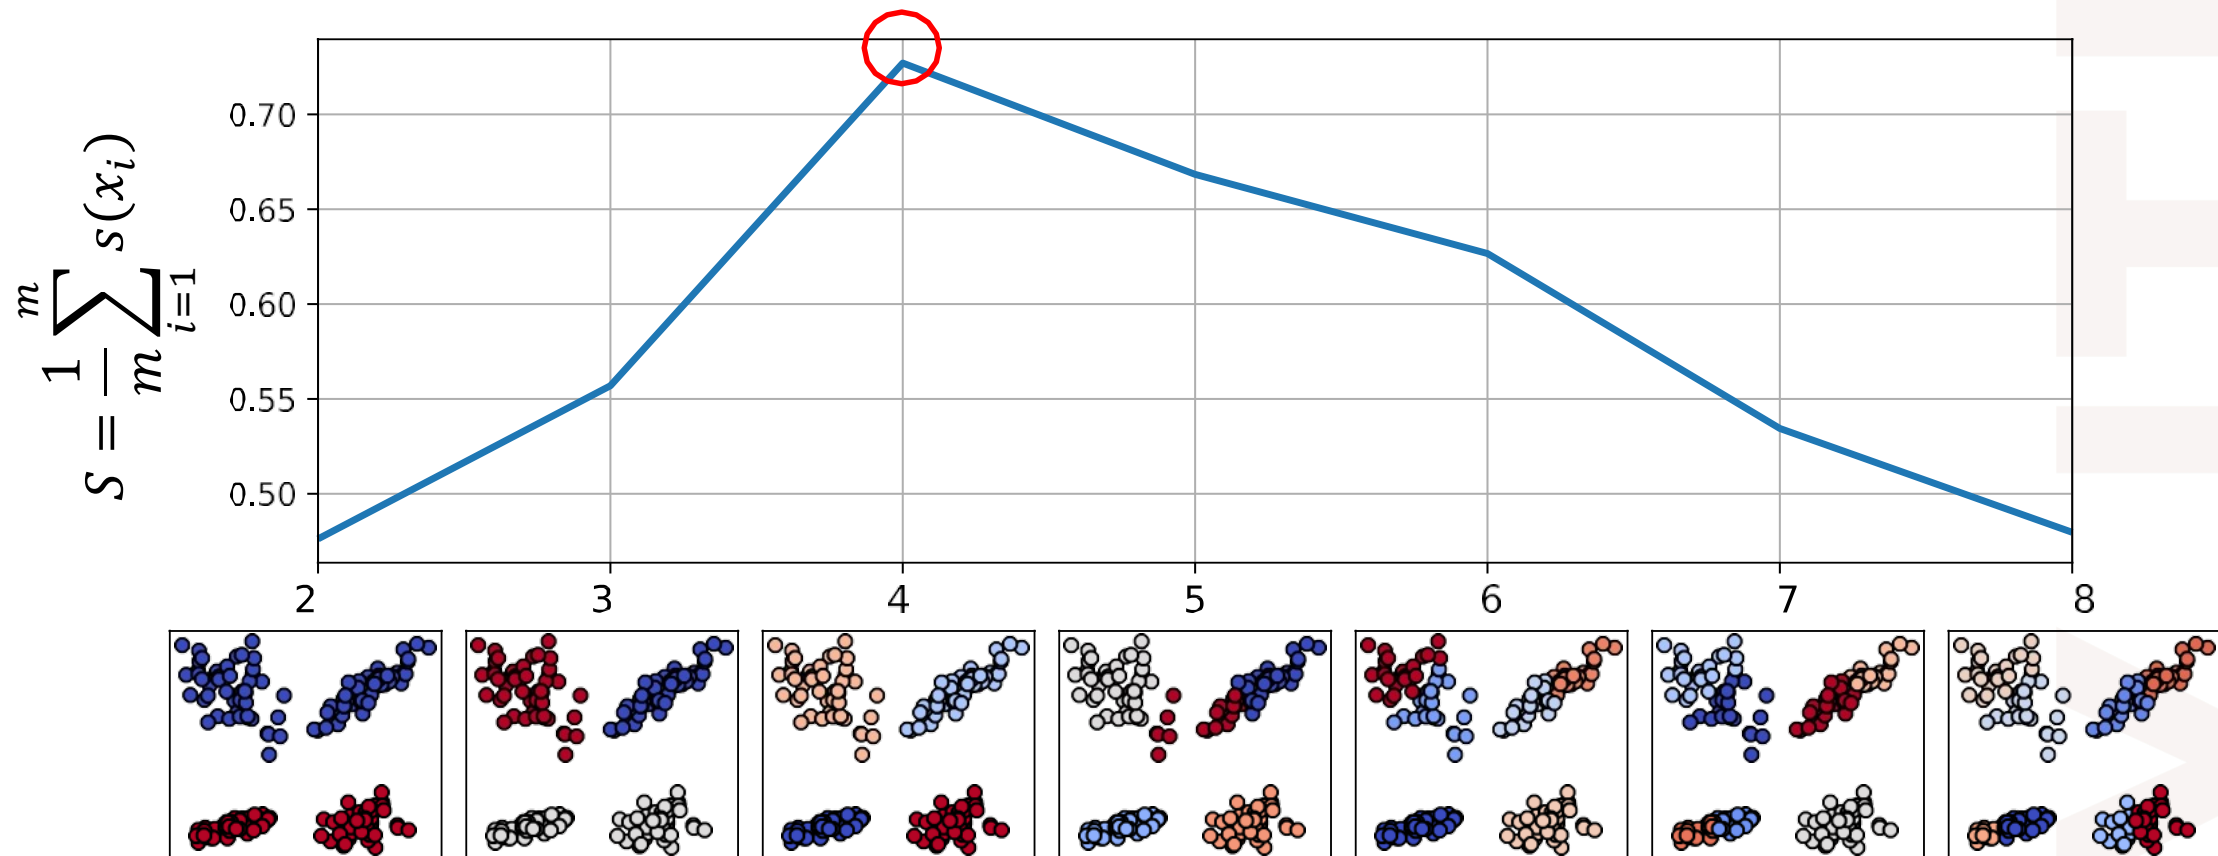
where:

- $a(x_i)$  is the average distance between  $x_i$  and all other points in the same cluster
- $b(x_i)$  is the average distance between  $x_i$  and all other points in the next neighbor cluster (i.e., the average distance to the nearest neighboring cluster)



# The Silhouette Coefficient


- Choose  $k$  that gives the highest mean silhouette



# K-means: pros and cons

- Pros:
  - Easy to implement
  - Scales to very large datasets
- Cons:
  - Difficult to choose  $K$ .
  - Only works on spherical, convex clusters.

Where K-means does not work well



# Hierarchical clustering

- Hierarchical Clustering is a set of clustering methods that aim at building a hierarchy of clusters
  - A cluster is composed of smaller clusters
- There are two strategies for building the hierarchy of clusters:
  - Agglomerative (bottom-up): we start with each point in its own cluster and we merge pairs of clusters until only one cluster is formed.
  - Divisive (top-down): we start with a single cluster containing the entire set of points and we recursively split until each point is in its own cluster.
- The most popular strategy in practical use is bottom-up (agglomerative)!



# Hierarchical clustering- Agglomerative


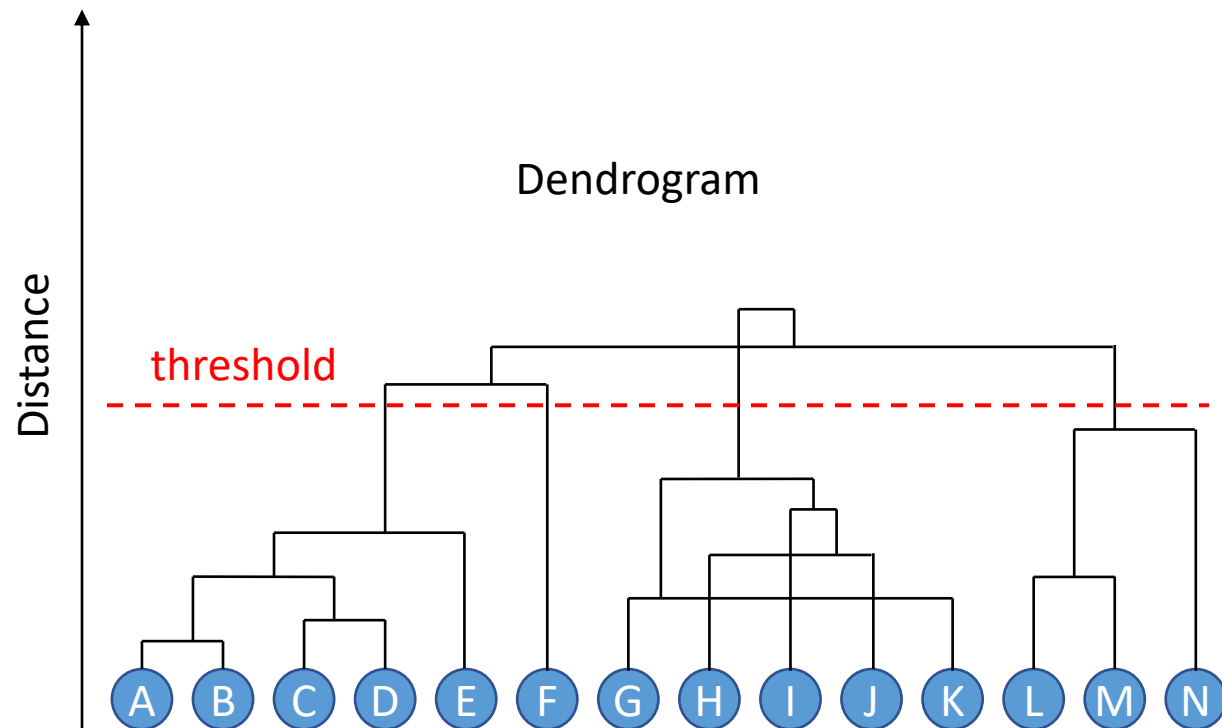
**Idea:** make sure *nearby* data points end up in the same cluster

- Initialize a collection  $\mathcal{C}$  of  $m$  singleton clusters, i.e.,  $c^{(i)} = \{x^{(i)}\}$
- Repeat until only one cluster is left:
  - Find a pair of clusters that is closest:  $\arg \min_{i,j} D(c^{(i)}, c^{(j)})$
  - Merge the two clusters  $c^{(i)}, c^{(j)}$  into a new cluster  $c^{(i \& j)}$
  - Remove  $c^{(i)}, c^{(j)}$  from the collection  $\mathcal{C}$ , and add  $c^{(i \& j)}$
- Produce a **dendrogram**: a hierarchical **tree** of clusters

Need to define **distance**




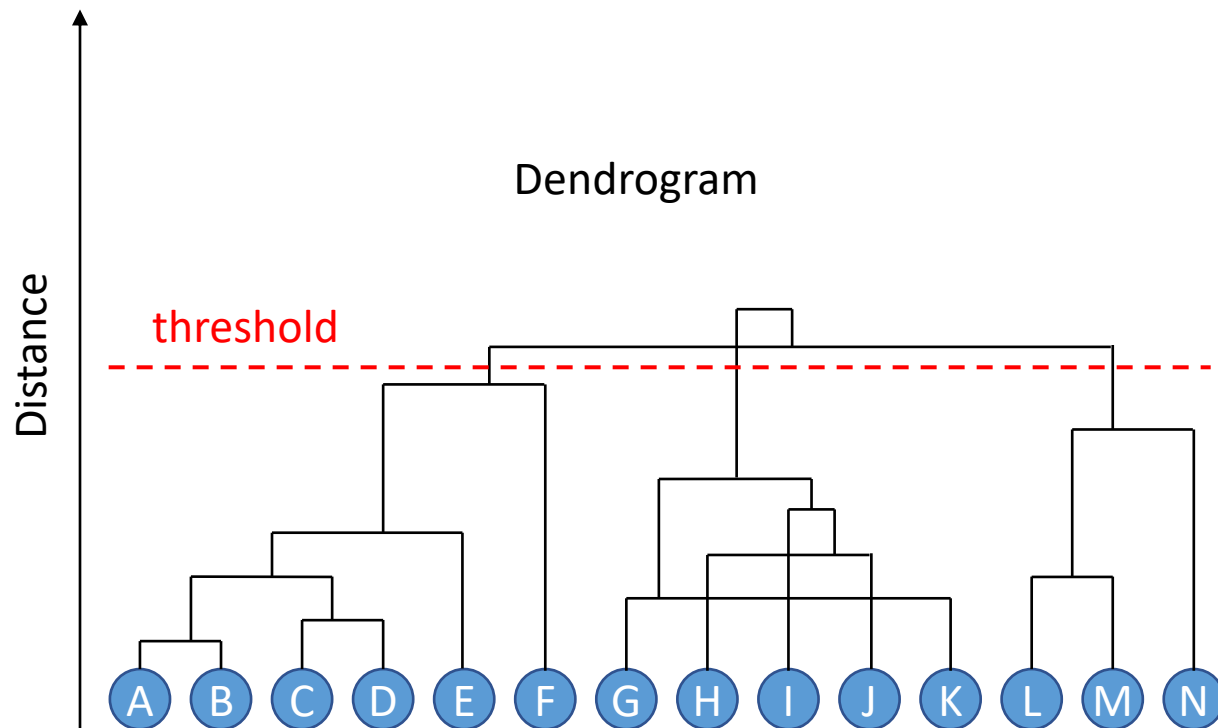
# Hierarchical clustering example





# Hierarchical clustering example


In this case, distance between clusters is defined by the **closest** pair



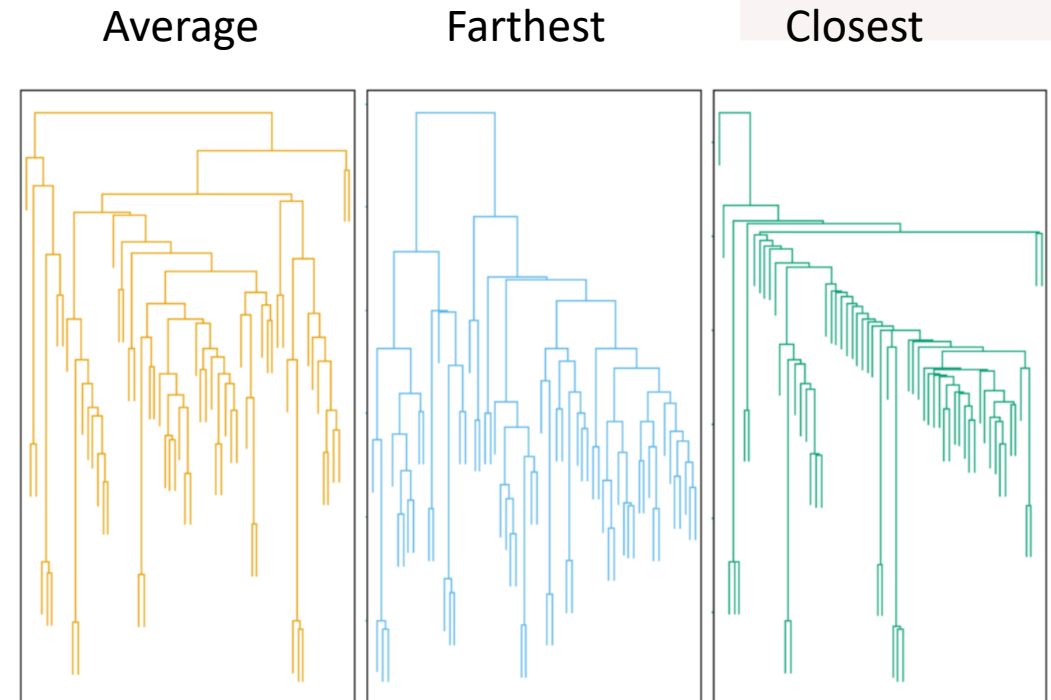
# Hierarchical clustering

Distance options:

- **single linkage (closest pair):**  
the minimum distance between samples in sub-clusters
- **complete linkage (farthest pair):**  
maximum distance between samples in sub-clusters
- **average linkage (average of all pairs):**  
average distance between each pair of samples in sub-clusters
- There are also other grouping strategies (such as centroid linkage)



Distance option influences the clustering result



# Hierarchical clustering

## Exercise

Given the following table that shows the distance between samples (“city block distance”), using agglomerative clustering method with **single linkage**, draw the final dendrogram obtained.

*city block distance :*

$$A(3,5) \ B(2,7) \rightarrow D(A,B) = |3-2| + |5-7| = 3$$


	A	B	C	D	E
A	0				
B	8	0			
C	3	6	0		
D	5	5	8	0	
E	13	10	2	7	0



# Hierarchical clustering: pros and cons

- Pros
  - Hierarchical structure is more informative than flat clusters ( $K$ -means)
  - Easier to decide the number of clusters
- Cons:
  - Slow to compute. Time complexity  $O(n^3)$ .
  - Sensitive to outliers, because it tries to connect all data points.

Hierarchical clustering with Ward's method



# Density-based clustering

**Idea:** Clustering based on density (local clustering criterion), e.g., number of *densely* connected points.


Most well-known algorithm: DBSCAN (Density-based spatial clustering of applications with noise )



# DBSCAN


Ester, et al (1996)

- DBSCAN classifies all points as core points, (*density*-)reachable points and outliers (or noise points):
- A point  $p$  is a core point if at least **minPts** points are within distance  $\epsilon$  ( $\epsilon$  is the maximum radius)
- A point  $q$  is directly reachable from  $p$  if point  $q$  is within distance  $\epsilon$  from point  $p$  and  $p$  must be a core point.
- A point  $q$  is reachable from  $p$  if there is a path  $p_1, \dots, p_n$  with  $p_1 = p$  and  $p_n = q$ , where each  $p_{i+1}$  is directly reachable from  $p_i$ .
- All points not reachable from any other point are outliers.



# DBSCAN breakdown


minPts = 4



Start from  $p_1$


$p_1$  is a core point. Create a new **cluster C1**

There are 4 neighbor points and they all become candidates to expand



Add  $p_2$  to C1

Found a new candidate  $p_3$




Add  $p_3$  to C1

Found a new neighbor  $p_4$ , but it cannot be a candidate


Because  $p_3$  is not a core point.




## DBSCAN breakdown (cont.)



Add  $p_5$  to C1  
No new candidate is found



Add  $p_6$  to C1  
A new candidate  $p_7$  is found




Add  $p_8$  to C1  
No new candidate is found







## DBSCAN breakdown (cont.)



Add  $p_9$  to C1  
A new candidate  $p_{10}$  is found



Add  $p_{10}$  to C1  
 $p_{10}$  is not core point, stop expanding




Mark noise points



## DBSCAN breakdown (cont.)

Final output



# DBSCAN pros and cons

- Pros:

- Pretty fast. Time complexity is  $O(n \log n)$  when optimized.
- Can find arbitrarily shaped clusters
- Robust to outliers (recognized as noise points)

- Cons:

- Cannot work well if density varies in different regions of data
- Choosing a proper distance threshold  $\epsilon$  can be difficult

